

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

## Exploratory Data Analysis

```
In [2]: df=pd.read_csv("cars.csv")
df
```

Out[2]:

	symboling	normalized- losses	make	fuel- type	body- style	drive- wheels	engine- location	width	height	en
0	3	?	alfa-romero	gas	convertible	rwd	front	64.1	48.8	
1	3	?	alfa-romero	gas	convertible	rwd	front	64.1	48.8	
2	1	?	alfa-romero	gas	hatchback	rwd	front	65.5	52.4	
3	2	164	audi	gas	sedan	fwd	front	66.2	54.3	
4	2	164	audi	gas	sedan	4wd	front	66.4	54.3	
...	...	...	...	...	...	...	...	...	...	
200	-1	95	volvo	gas	sedan	rwd	front	68.9	55.5	
201	-1	95	volvo	gas	sedan	rwd	front	68.8	55.5	
202	-1	95	volvo	gas	sedan	rwd	front	68.9	55.5	
203	-1	95	volvo	diesel	sedan	rwd	front	68.9	55.5	
204	-1	95	volvo	gas	sedan	rwd	front	68.9	55.5	

205 rows × 15 columns



```
In [3]: df.shape
```

Out[3]: (205, 15)

```
In [4]: df.describe()
```

Out[4]:

	symboling	width	height	engine-size	city-mpg	highway-mpg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	0.834146	65.907805	53.724878	126.907317	25.219512	30.751220	13227.4780
std	1.245307	2.145204	2.443522	41.642693	6.542142	6.886443	7902.6516
min	-2.000000	60.300000	47.800000	61.000000	13.000000	16.000000	5118.0000
25%	0.000000	64.100000	52.000000	97.000000	19.000000	25.000000	7788.0000
50%	1.000000	65.500000	54.100000	120.000000	24.000000	30.000000	10345.0000
75%	2.000000	66.900000	55.500000	141.000000	30.000000	34.000000	16500.0000
max	3.000000	72.300000	59.800000	326.000000	49.000000	54.000000	45400.0000

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    int64
1   normalized-losses      205 non-null    object
2   make                   205 non-null    object
3   fuel-type              205 non-null    object
4   body-style             205 non-null    object
5   drive-wheels           205 non-null    object
6   engine-location        205 non-null    object
7   width                  205 non-null    float64
8   height                 205 non-null    float64
9   engine-type            205 non-null    object
10  engine-size            205 non-null    int64
11  horsepower              205 non-null    object
12  city-mpg                205 non-null    int64
13  highway-mpg            205 non-null    int64
14  price                  205 non-null    int64
dtypes: float64(2), int64(5), object(8)
memory usage: 24.1+ KB
```

## Handling Missing Values

```
In [6]: df.isna()
```

Out[6]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
0	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...	...	...	...
200	False	False	False	False	False	False	False	False	False	False
201	False	False	False	False	False	False	False	False	False	False
202	False	False	False	False	False	False	False	False	False	False
203	False	False	False	False	False	False	False	False	False	False
204	False	False	False	False	False	False	False	False	False	False

205 rows × 15 columns



```
In [7]: df.isna().sum()
```

Out[7]:

symboling	0
normalized-losses	0
make	0
fuel-type	0
body-style	0
drive-wheels	0
engine-location	0
width	0
height	0
engine-type	0
engine-size	0
horsepower	0
city-mpg	0
highway-mpg	0
price	0
dtype:	int64

## Changing Dtype from Object to Float64

```
In [8]: >>> '''df["normalized-losses"].replace("?", np.nan, inplace=True)
df["horsepower"].replace("?", np.nan, inplace=True)

df["normalized-losses"] = df["normalized-losses"].astype("float64")
df["horsepower"] = df["horsepower"].astype("float64")

nmean = df["normalized-losses"].mean()
hmean = df["horsepower"].mean()

df["normalized-losses"].fillna(nmean, inplace=True)
df["horsepower"].fillna(hmean, inplace=True)'''
```

```
Out[8]: 'df["normalized-losses"].replace("?", np.nan, inplace=True)\ndf["horsepower"].replace("?", np.nan, inplace=True)\n\ndf["normalized-losses"] = df["normalized-losses"].astype("float64")\ndf["horsepower"] = df["horsepower"].astype("float64")\n\nnmean = df["normalized-losses"].mean()\nhmean = df["horsepower"].mean()\n\ndf["normalized-losses"].fillna(nmean, inplace=True)\ndf["horsepower"].fillna(hmean, inplace=True)'
```

```
In [9]: >>> df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    int64
1   normalized-losses      205 non-null    object
2   make                   205 non-null    object
3   fuel-type              205 non-null    object
4   body-style             205 non-null    object
5   drive-wheels           205 non-null    object
6   engine-location        205 non-null    object
7   width                  205 non-null    float64
8   height                 205 non-null    float64
9   engine-type            205 non-null    object
10  engine-size            205 non-null    int64
11  horsepower             205 non-null    object
12  city-mpg               205 non-null    int64
13  highway-mpg            205 non-null    int64
14  price                  205 non-null    int64
dtypes: float64(2), int64(5), object(8)
memory usage: 24.1+ KB
```

```
In [10]: >>> #df["normalized-losses"]
```

```
In [11]: >>> #df["horsepower"]
```

```
In [12]: df["normalized-losses"].replace("?", np.nan, inplace=True)
df["horsepower"].replace("?", np.nan, inplace=True)

df["normalized-losses"] = df["normalized-losses"].astype("float64")
df["horsepower"] = df["horsepower"].astype("float64")
```

```
In [13]: from sklearn.impute import SimpleImputer
```

```
In [14]: si = SimpleImputer(missing_values=np.nan, strategy="mean")
```

```
In [15]: df[["normalized-losses", "horsepower"]] = si.fit_transform(df[["normalized-lo
```

```
In [16]: df["normalized-losses"]
```

```
Out[16]: 0      122.0
1      122.0
2      122.0
3      164.0
4      164.0
...
200     95.0
201     95.0
202     95.0
203     95.0
204     95.0
Name: normalized-losses, Length: 205, dtype: float64
```

```
In [17]: df["horsepower"]
```

```
Out[17]: 0      111.0
1      111.0
2      154.0
3      102.0
4      115.0
...
200    114.0
201    160.0
202    134.0
203    106.0
204    114.0
Name: horsepower, Length: 205, dtype: float64
```

In [18]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   symboling              205 non-null   int64  
1   normalized-losses      205 non-null   float64
2   make                   205 non-null   object  
3   fuel-type              205 non-null   object  
4   body-style             205 non-null   object  
5   drive-wheels           205 non-null   object  
6   engine-location        205 non-null   object  
7   width                  205 non-null   float64
8   height                 205 non-null   float64
9   engine-type            205 non-null   object  
10  engine-size            205 non-null   int64  
11  horsepower             205 non-null   float64
12  city-mpg               205 non-null   int64  
13  highway-mpg            205 non-null   int64  
14  price                  205 non-null   int64  
dtypes: float64(4), int64(5), object(6)
memory usage: 24.1+ KB
```

In [19]: `df.head()`

Out[19]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
0	3	122.0	alfa-romero	gas	convertible	rwd	front	64.1	48.8	doh
1	3	122.0	alfa-romero	gas	convertible	rwd	front	64.1	48.8	doh
2	1	122.0	alfa-romero	gas	hatchback	rwd	front	65.5	52.4	ohc
3	2	164.0	audi	gas	sedan	fwd	front	66.2	54.3	oh
4	2	164.0	audi	gas	sedan	4wd	front	66.4	54.3	oh

```
In [20]: df.describe()
```

Out[20]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mp
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	0.834146	122.000000	65.907805	53.724878	126.907317	104.256158	25.21951
std	1.245307	31.681008	2.145204	2.443522	41.642693	39.519211	6.54214
min	-2.000000	65.000000	60.300000	47.800000	61.000000	48.000000	13.00000
25%	0.000000	101.000000	64.100000	52.000000	97.000000	70.000000	19.00000
50%	1.000000	122.000000	65.500000	54.100000	120.000000	95.000000	24.00000
75%	2.000000	137.000000	66.900000	55.500000	141.000000	116.000000	30.00000
max	3.000000	256.000000	72.300000	59.800000	326.000000	288.000000	49.00000

```
In [21]: feature=df.iloc[:, :-1]
target=df["price"]
```

```
In [22]: feature
```

Out[22]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	en
0	3	122.0	alfa-romero	gas	convertible	rwd	front	64.1	48.8	
1	3	122.0	alfa-romero	gas	convertible	rwd	front	64.1	48.8	
2	1	122.0	alfa-romero	gas	hatchback	rwd	front	65.5	52.4	
3	2	164.0	audi	gas	sedan	fwd	front	66.2	54.3	
4	2	164.0	audi	gas	sedan	4wd	front	66.4	54.3	
...	...	...	...	...	...	...	...	...	...	
200	-1	95.0	volvo	gas	sedan	rwd	front	68.9	55.5	
201	-1	95.0	volvo	gas	sedan	rwd	front	68.8	55.5	
202	-1	95.0	volvo	gas	sedan	rwd	front	68.9	55.5	
203	-1	95.0	volvo	diesel	sedan	rwd	front	68.9	55.5	
204	-1	95.0	volvo	gas	sedan	rwd	front	68.9	55.5	

205 rows × 14 columns


In [23]:  target

```
Out[23]: 0      13495
          1      16500
          2      16500
          3      13950
          4      17450
          ...
          200    16845
          201    19045
          202    21485
          203    22470
          204    22625
          Name: price, Length: 205, dtype: int64
```

## Handling outliers



```
In [24]: ▶ plt.figure(figsize=(20,20))
plt.grid()
sns.boxplot(data=feature,x=target);
```



```
In [25]: ▶ df[(df["make"]=="dodge")&(df["price"]>9000)]
```

Out[25]:


	symboling	normalized- losses	make	fuel- type	body- style	drive- wheels	engine- location	width	height	engine typ
29	3	145.0	dodge	gas	hatchback	fwd	front	66.3	50.2	oh

```
In [26]: ▶ df.loc[29,"price"]
```

Out[26]: 12964

```
In [27]: df.loc[29,"price"]=9000
```

```
In [28]: plt.figure(figsize=(20,20))
plt.grid()
sns.boxplot(data=feature,x=target,y="make");
```




```
In [29]: df[(df["make"]=="honda")&(df["price"]>12500)]
```

Out[29]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
41	0	85.0	honda	gas	sedan	fwd	front	65.2	54.1	ohc

```
In [30]: df.loc[41,"price"]=12000
```

```
In [31]: plt.figure(figsize=(20,20))
plt.grid()
sns.boxplot(data=feature,x=target,y="make");
```



```
In [32]: df[(df["make"]=="toyota")&(df["price"]>13500)]
```

Out[32]:


	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
172	2	134.0	toyota	gas	convertible	rwd	front	65.6	53.0	4-cyl
178	3	197.0	toyota	gas	hatchback	rwd	front	67.7	52.0	4-cyl
179	3	197.0	toyota	gas	hatchback	rwd	front	67.7	52.0	4-cyl
180	-1	90.0	toyota	gas	sedan	rwd	front	66.5	54.1	4-cyl
181	-1	122.0	toyota	gas	wagon	rwd	front	66.5	54.1	4-cyl

```
In [33]: df.loc[[172,178,179,180,181],"price"]
```

```
Out[33]: 172    17669
178    16558
179    15998
180    15690
181    15750
Name: price, dtype: int64
```

```
In [34]: df.loc[[172,178,179,180,181],"price"]=13000
```

```
In [35]: plt.figure(figsize=(20,20))
plt.grid()
sns.boxplot(data=feature,x=target,y="make");
```



```
In [36]: df[(df["make"]=="mitsubishi")&(df["price"]>13000)]
```

Out[36]:

	symboling	normalized- losses	make	fuel- type	body- style	drive- wheels	engine- location	width	height	en
83	3	122.0	mitsubishi	gas	hatchback	fwd	front	66.3	50.2	
84	3	122.0	mitsubishi	gas	hatchback	fwd	front	66.3	50.2	




```
In [37]: df.loc[[83,84],"price"]
```

Out[37]: 83 14869  
84 14489  
Name: price, dtype: int64

```
In [38]: df.loc[[83,84],"price"]=13000
```

```
In [39]: plt.figure(figsize=(20,20))
plt.grid()
sns.boxplot(data=feature,x=target,y="make");
```



```
In [40]: df[(df["make"]=="plymouth")&(df["price"]>10000)]
```

Out[40]:


	symboling	normalized- losses	make	fuel- type	body- style	drive- wheels	engine- location	width	height	en
124	3	122.0	plymouth	gas	hatchback	rwd	front	66.3	50.2	

```
In [41]: df.loc[124,"price"]
```

Out[41]: 12764

```
In [42]: df.loc[124,"price"]=8500
```

```
In [43]: plt.figure(figsize=(20,20))
plt.grid()
sns.boxplot(data=feature,x=target,y="make");
```



```
In [44]: df[(df["make"]=="isuzu")&(df["price"]>22500)]
```

Out[44]:


	symboling	normalized- losses	make	fuel- type	body- style	drive- wheels	engine- location	width	height	engine- type	e
45	0	122.0	isuzu	gas	sedan	fwd	front	63.6	52.0	ohc	

```
In [45]: df.loc[45,"price"]
```

Out[45]: 23000

```
In [46]: df.loc[45, "price"] = 13500
```

```
In [47]: plt.figure(figsize=(20,20))  
plt.grid()  
sns.boxplot(data=feature, x=target, y="make");
```



**To identify skew from numeric datatype and how to apply log transformation**



```
In [48]: df.describe()
```

Out[48]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	0.834146	122.000000	65.907805	53.724878	126.907317	104.256158	25.21951
std	1.245307	31.681008	2.145204	2.443522	41.642693	39.519211	6.54214
min	-2.000000	65.000000	60.300000	47.800000	61.000000	48.000000	13.00000
25%	0.000000	101.000000	64.100000	52.000000	97.000000	70.000000	19.00000
50%	1.000000	122.000000	65.500000	54.100000	120.000000	95.000000	24.00000
75%	2.000000	137.000000	66.900000	55.500000	141.000000	116.000000	30.00000
max	3.000000	256.000000	72.300000	59.800000	326.000000	288.000000	49.00000

```
In [49]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    int64
1   normalized-losses      205 non-null    float64
2   make                   205 non-null    object
3   fuel-type              205 non-null    object
4   body-style             205 non-null    object
5   drive-wheels           205 non-null    object
6   engine-location        205 non-null    object
7   width                  205 non-null    float64
8   height                 205 non-null    float64
9   engine-type            205 non-null    object
10  engine-size            205 non-null    int64
11  horsepower              205 non-null    float64
12  city-mpg               205 non-null    int64
13  highway-mpg            205 non-null    int64
14  price                  205 non-null    int64
dtypes: float64(4), int64(5), object(6)
memory usage: 24.1+ KB
```

```
In [50]: feature.select_dtypes(["int64", "float64"]).columns
```

Out[50]: Index(['symboling', 'normalized-losses', 'width', 'height', 'engine-size',  
'horsepower', 'city-mpg', 'highway-mpg'],  
dtype='object')

```
In [51]: colname=feature.select_dtypes(["int64","float64"]).columns
```

```
In [52]: colname
```

```
Out[52]: Index(['symboling', 'normalized-losses', 'width', 'height', 'engine-size',  
              'horsepower', 'city-mpg', 'highway-mpg'],  
              dtype='object')
```

```
In [53]: from scipy.stats import skew
```

```
In [54]: skew(feature["normalized-losses"])
```

```
Out[54]: 0.8485348696008058
```

```
In [55]: feature[colname]
```


	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg	highway-mpg
0	3	122.0	64.1	48.8	130	111.0	21	27
1	3	122.0	64.1	48.8	130	111.0	21	27
2	1	122.0	65.5	52.4	152	154.0	19	26
3	2	164.0	66.2	54.3	109	102.0	24	30
4	2	164.0	66.4	54.3	136	115.0	18	22
...	...	...	...	...	...	...	...	...
200	-1	95.0	68.9	55.5	141	114.0	23	28
201	-1	95.0	68.8	55.5	141	160.0	19	25
202	-1	95.0	68.9	55.5	173	134.0	18	23
203	-1	95.0	68.9	55.5	145	106.0	26	27
204	-1	95.0	68.9	55.5	141	114.0	19	25

205 rows x 9 columns

```
In [56]: skew(feature[colname])
```

```
Out[56]: array([0.20952469, 0.84853487, 0.89737535, 0.06265992, 1.93337485,  
               1.38751473, 0.65883775, 0.53603793])
```

```
In [57]: for i in feature[colname]:
          print(i)
          print(skew(feature[i]))
          plt.figure()
          sns.distplot(feature[i])
          plt.show()
```



```
In [58]: df.corr()# to check the corelation
```

C:\Users\Reshmi\AppData\Local\Temp\ipykernel\_21608\3597234383.py:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
df.corr()# to check the corelation
```

Out[58]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg
symboling	1.000000	0.465190	-0.232919	-0.541038	-0.105790	0.071389	-0.035823
normalized-losses	0.465190	1.000000	0.084195	-0.370706	0.110997	0.203434	-0.218749
width	-0.232919	0.084195	1.000000	0.279210	0.735433	0.642195	-0.642704
height	-0.541038	-0.370706	0.279210	1.000000	0.067149	-0.110137	-0.048640
engine-size	-0.105790	0.110997	0.735433	0.067149	1.000000	0.810713	-0.653658
horsepower	0.071389	0.203434	0.642195	-0.110137	0.810713	1.000000	-0.803162
city-mpg	-0.035823	-0.218749	-0.642704	-0.048640	-0.653658	-0.803162	1.000000
highway-mpg	0.034606	-0.178221	-0.677218	-0.107358	-0.677470	-0.770903	0.971337
price	-0.099208	0.123851	0.722863	0.150782	0.848517	0.736585	-0.655155

```
In [59]: ▶ pd.concat([feature,target],axis=1).corr().style.background_gradient()
```

C:\Users\Reshmi\AppData\Local\Temp\ipykernel\_21608\4285486589.py:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
pd.concat([feature,target],axis=1).corr().style.background_gradient()
```

Out[59]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg
symboling	1.000000	0.465190	-0.232919	-0.541038	-0.105790	0.071389	-0.035823
normalized-losses	0.465190	1.000000	0.084195	-0.370706	0.110997	0.203434	-0.218749
width	-0.232919	0.084195	1.000000	0.279210	0.735433	0.642195	-0.642704
height	-0.541038	-0.370706	0.279210	1.000000	0.067149	-0.110137	-0.048640
engine-size	-0.105790	0.110997	0.735433	0.067149	1.000000	0.810713	-0.653658
horsepower	0.071389	0.203434	0.642195	-0.110137	0.810713	1.000000	-0.803162
city-mpg	-0.035823	-0.218749	-0.642704	-0.048640	-0.653658	-0.803162	1.000000
highway-mpg	0.034606	-0.178221	-0.677218	-0.107358	-0.677470	-0.770903	0.971337
price	-0.099208	0.123851	0.722863	0.150782	0.848517	0.736585	-0.655155



```
In [60]: ▶ df["normalized-losses"].unique()
```

Out[60]: array([122., 164., 158., 192., 188., 121., 98., 81., 118., 148., 110., 145., 137., 101., 78., 106., 85., 107., 104., 113., 150., 129., 115., 93., 142., 161., 153., 125., 128., 103., 168., 108., 194., 231., 119., 154., 74., 186., 83., 102., 89., 87., 77., 91., 134., 65., 197., 90., 94., 256., 95.] )

```
In [61]: ▶ feature["normalized-losses"]=np.log(feature["normalized-losses"])
```

```
In [62]: ▶ skew(feature["normalized-losses"])
```

Out[62]: 0.03137735337911685

## Encoding

1.OneHotEncoding 2.LabelEncoding

```
In [63]: ▶ from sklearn.preprocessing import OneHotEncoder
```

```
In [64]: one=OneHotEncoder()  
one.fit_transform(feature[["make"]]).toarray()
```

```
Out[64]: array([[1., 0., 0., ..., 0., 0., 0.],  
                [1., 0., 0., ..., 0., 0., 0.],  
                [1., 0., 0., ..., 0., 0., 0.],  
                ...,  
                [0., 0., 0., ..., 0., 0., 1.],  
                [0., 0., 0., ..., 0., 0., 1.],  
                [0., 0., 0., ..., 0., 0., 1.]])
```

```
In [65]: from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()  
le.fit_transform(target)
```

```
Out[65]: array([119, 134, 134, 125, 144, 128, 145, 153, 163, 109, 133, 141, 157,  
                158, 164, 168, 180, 177,  1,  11,  17,  7,  13,  49,  10,  20,  
                41,  64,  67,  70,  14,  25,  5,  16,  32,  34,  34,  47,  71,  
                66,  91, 109,  92,  22,  88, 121,  98, 170, 175, 176,  2,  8,  
                23,  21,  37,  97, 106, 122, 131,  66,  61,  93,  90,  95, 100,  
                148, 149, 165, 167, 166, 169, 173, 174, 179, 181, 135,  4,  9,  
                19,  42,  82,  62, 113, 115, 115, 28,  55,  74,  74,  6,  30,  
                18,  24,  36,  35,  46,  39,  52,  58,  69,  79, 120, 126, 120,  
                143, 156, 150, 108, 116, 112, 124, 130, 140, 138, 142, 137, 146,  
                147,  7,  49,  10,  20,  41,  67,  63, 160, 171, 172, 178, 109,  
                75,  81, 107, 110, 127, 129, 147, 152,  0,  29,  40,  31,  44,  
                83,  72, 102,  38,  89,  53, 105,  3,  12,  15,  26,  48,  65,  
                27,  33,  48,  45,  43,  59,  73,  54,  57,  76,  78,  60,  80,  
                86,  99, 103, 115,  68,  94,  85,  96, 101, 115, 115, 115, 115,  
                44,  50,  51,  56,  61,  77,  87, 104,  84, 117, 123, 111, 114,  
                118, 132, 136, 151, 154, 139, 155, 159, 161, 162], dtype=int64)
```

```
In [66]: catcol=feature.select_dtypes("object").columns
```

```
In [67]: catcol
```

```
Out[67]: Index(['make', 'fuel-type', 'body-style', 'drive-wheels', 'engine-locatio  
n',  
                'engine-type'],  
               dtype='object')
```

## Using OrdinalEncoder

input

```
In [68]: from sklearn.preprocessing import OrdinalEncoder
         oe=OrdinalEncoder()
         feature[catcol]=oe.fit_transform(feature[catcol])
```

```
In [69]: feature[catcol]
```

Out[69]:

	make	fuel-type	body-style	drive-wheels	engine-location	engine-type
0	0.0	1.0	0.0	2.0	0.0	0.0
1	0.0	1.0	0.0	2.0	0.0	0.0
2	0.0	1.0	2.0	2.0	0.0	5.0
3	1.0	1.0	3.0	1.0	0.0	3.0
4	1.0	1.0	3.0	0.0	0.0	3.0
...	...	...	...	...	...	...
200	21.0	1.0	3.0	2.0	0.0	3.0
201	21.0	1.0	3.0	2.0	0.0	3.0
202	21.0	1.0	3.0	2.0	0.0	5.0
203	21.0	0.0	3.0	2.0	0.0	3.0
204	21.0	1.0	3.0	2.0	0.0	3.0

205 rows × 6 columns

```
In [70]: feature.head()
```

Out[70]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
0	3	4.804021	0.0	1.0	0.0	2.0	0.0	64.1	48.8	0.0
1	3	4.804021	0.0	1.0	0.0	2.0	0.0	64.1	48.8	0.0
2	1	4.804021	0.0	1.0	2.0	2.0	0.0	65.5	52.4	5.0
3	2	5.099866	1.0	1.0	3.0	1.0	0.0	66.2	54.3	3.0
4	2	5.099866	1.0	1.0	3.0	0.0	0.0	66.4	54.3	3.0



```
In [ ]:
```